

Untyped Allocation

Dan Pittman dan@auxon.io

September 10, 2021

Introduction

In seL4 parlance, an **Untyped** is a kernel object representing a yet-to-be-spoken-for chunk of physical memory. When a kernel object is created, it must be done so through the seL4 system call `seL4_Untyped_Retype`. `seL4_Untyped_Retype` take an untyped kernel object, an object representing the kernel object to retype this untyped chunk into, as well as the target object's *size*. That size is used by the kernel to track offsets in an untyped object. Sizes are described in terms of the number of "bits" in the object's size, where a bit x is used in the following equation $2^x = \text{Size}$. For example, a object whose size is 4, can be thought of as 2^4 bytes, or 16 bytes.

Because **Untyped**s are themselves kernel objects with an associated, but variable, size, they can also be the target of a `seL4_Untyped_Retype` invocation. In Ferros this is common practice: A process acquires an **Untyped** large enough to hold all of the capabilities it intends to create, then breaks that object down into the necessary sizes in a \log_2 fashion. I.e., an untyped object of size 6 can be broken into 2 of size 5, 4 of size 4, 8 of size 3, and so on.

Buddy Allocator

A *buddy* allocator is one which does this \log_2 break down automatically by being given a desired allocation size, and then recursively breaking down the object until an object of that size is available. In Ferros, a buddy allocator is used by wrapping an **Untyped**, and then using the buddy algorithm to allocate chunks of it.

Construction

`OneHotUList` builds a `UList`, the buddy allocator's state, given the index at which that initial object ought to live. The size, sz , of the objects in the list are determined via the equation $sz = 2^{i+\text{MinUtSize}}$.

`OneHotUList`'s definition:

```

pub trait _OneHotUList: Unsigned {
  type Output;
}

type OneHotUList<Index> = <Index as _OneHotUList>::Output;

impl _OneHotUList for U0 {
  type Output = ULCons<U1, UNull>;
}

impl<IHead: Bit, ITail: Unsigned> _OneHotUList for UInt<ITail, IHead>
where
  UInt<ITail, IHead>: Sub<U1>,
  Diff<UInt<ITail, IHead>, U1>: _OneHotUList,
  OneHotUList<Diff<UInt<ITail, IHead>, U1>>: UList,
{
  type Output = ULCons<U0, OneHotUList<Diff<UInt<ITail, IHead>, U1>>>;
}

```

Is directly translated to the `buddy` constructor:

```

buddy : ∀ n → Vec ℕ (suc n)
buddy zero = 1 :: []
buddy (suc n) = 0 :: buddy n

```

Allocation

To allocate an `Untyped`, we take its expected position in the list which can be acquired by solving for i in the equation stated in the previous section. Once we have the index, we begin the process of folding the buddy allocator's state, a list of available `Untyped`s, splitting if necessary until we have the size we'd set out for. `get-untyped` also tracks its number of splits. This is because each split creates a new kernel object whose `cptr` must have a slot to live in; the split count tells us how many cspace slots will be needed.

$$\text{get-untyped} : \forall \{n\} \rightarrow \text{Fin } n \rightarrow \text{Vec } \mathbb{N} \ n \rightarrow \text{Maybe } (\mathbb{N} \times \text{Vec } \mathbb{N} \ n)$$

As stated before `get-untyped` is implemented as a fold over the allocator's state.

```

get-untyped index untyped =
  let state = foldr GetUtState
    fold-ut
    record { uts = []
      ; idx = (toN index)
      ; splits = 0
      ; done = false
    }
    untyped
  in if (done state) then
    just ((splits state) , (uts state))
  else nothing

```

The fold function uses with-abstraction to build a truth table telling us whether:

1. We've found the object we're looking for.
2. Whether the index we've been given is zero.
3. Whether the untyped in the current position is absent.

```

fold-ut : ∀ {n} → ℕ → GetUtState n → GetUtState (suc n)
fold-ut ut state with done state | is-zero (idx state) | is-zero ut

```

We will break down the meaning of the truth table's states one by one.

1. T | _ | _

We've found the untyped we're looking for and can be done. We just tack on the remaining untyped as they are.

```

... | true | _ | _ = record { uts = ut :: (uts state)
  ; idx = (idx state)
  ; splits = (splits state)
  ; done = (done state)
}

```

2. F | F | _

We haven't reached our desired untyped size yet, our index is non-zero.

```

... | false | false | _ = record { uts = ut :: (uts state)
  ; idx = pred (idx state)
  ; splits = (splits state)
  ; done = (done state)
}

```

3. F | T | F

The index is zero and this cell has an untyped for us to take. Take it, mark it as done, move on.

```
... | false | true | false = record { uts = (pred ut) :: (uts state)
                                   ; idx = (idx state)
                                   ; splits = (splits state)
                                   ; done = true
                                   }
```

4. F | T | T

We've reached our index (or moved past it), and have not yet found an untyped of the right size. So we add one untyped in this cell, and count our splits.

```
... | false | true | true = record { uts = (suc ut) :: (uts state)
                                   ; idx = (idx state)
                                   ; splits = suc (splits state)
                                   ; done = (done state)
                                   }
```

This code corresponds to the following type-level implementation in Rust:

```
/// Type-level function to track the result of an allocation
pub trait _TakeUntyped<Index> {
    type ResultPoolSizes;
    type NumSplits;
}

// Index is non-zero, and there are pools left: recur with Index-1, and the
// remaining pools
impl<IndexU: Unsigned, IndexB: Bit, Head: Unsigned, Tail: UList>
    _TakeUntyped<UInt<IndexU, IndexB>>
    for ULCons<Head, Tail>
where
    UInt<IndexU, IndexB>: Sub<U1>,
    Diff<UInt<IndexU, IndexB>, U1>: Unsigned,

    Tail: _TakeUntyped<Diff<UInt<IndexU, IndexB>, U1>>,
    TakeUntyped_ResultPoolSizes<Tail, Diff<UInt<IndexU, IndexB>, U1>>: UList,
{
    type ResultPoolSizes =
        ULCons<Head,
```

```

        TakeUntyped_ResultPoolSizes<Tail,
                                Diff<UInt<IndexU, IndexB>, U1>>>;
    type NumSplits = TakeUntyped_NumSplits<Tail,
                                Diff<UInt<IndexU, IndexB>, U1>>>;
}

// Index is 0, and the head pool has resources: remove one from it,
// with no splits.
impl<HeadU: Unsigned, HeadB: Bit, Tail: UList>
    _TakeUntyped<U0> for ULCons<UInt<HeadU, HeadB>, Tail>
where
    UInt<HeadU, HeadB>: Sub<U1>,
    Diff<UInt<HeadU, HeadB>, U1>: Unsigned,
{
    type ResultPoolSizes = ULCons<Diff<UInt<HeadU, HeadB>, U1>, Tail>;
    type NumSplits = U0;
}

// index is zero, and the head pool is empty. Take one from the next
// pool (which we will split, and return one of), and put one (the
// remainder) in the head pool.
impl<Tail: UList> _TakeUntyped<U0> for ULCons<U0, Tail>
where
    Tail: _TakeUntyped<U0>,
    U1: Add<TakeUntyped_NumSplits<Tail, U0>>,
    Sum<U1, TakeUntyped_NumSplits<Tail, U0>>: Unsigned,
    TakeUntyped_ResultPoolSizes<Tail, U0>: UList,
{
    type ResultPoolSizes = ULCons<U1, TakeUntyped_ResultPoolSizes<Tail, U0>>;
    type NumSplits = Sum<U1, TakeUntyped_NumSplits<Tail, U0>>;
}

```

The differences to note are that, of course, Rust's type system does not provide a right fold function. Instead, the truth table mentioned above is broken into the "match clauses"—the different type-level patterns for which `_TakeUntyped` is implemented.

Secondarily, in Agda, our implementation must be total, this is what we get through our use of `Fin`, for instance—an index cannot be larger than the vector comprising the allocator's state. However in Rust, we intentionally write partial type-level functions leaving the unhandled cases to cause a compilation error. For example, in Rust, if a request comes in for an unavailable size, we leave that case dangling, however in Agda, we must handle it via a `Maybe`.